

Comparative Analysis of Machine Learning and Deep Learning Algorithms for Spam Detection

Srivatsa Chidara

s.chidara@student.utwente.nl

s3105520

Abstract

This project investigates spam classification performance by comparing various machine learning and deep learning models, this include Support Vector Machine(SVM), Random Forest(RF), Bidirectional Long Short-Term Memory RNN(BiLSTM), and ALBERT. The analysis examines the effects of two feature extraction methods, namely Term Frequency - Inverse Document Frequency (TF-IDF) and Word2Vec on machine learning algorithms. Various pre-processing steps are examined to determine their effect on model performance. Results indicate that SVM with TF-IDF performs best. Pre-processing contributes positively to most models, but shows minimal impact on ALBERT. This project provide insights into feature selection and role of pre-processing, which can be used in research of robust spam detection.

1 Introduction

In today's digitally connected world, the amount of information exchanged through electronic communication channels, such as email, messaging applications, and social media has grown exponentially. Sadly, this growth has been accompanied by a significance rise in potentially harmful messages which are called as spam. Spam messages are not just junk, and clutter in the inbox. They serve as the gateway for many malicious activities such as phishing, malware, and various other kinds of malicious activities which can compromise the security. This project aims to explore various methods with a goal of improving the accuracy for the spam classification. Specially, I focus on comparing the performance among 4 algorithms,machine learning algorithms such as Random Forest, and Support Vector Machine, alongside the deep learning models, such Bidirectional Long Short-Term Memory (BiLSTM) (**BILSTM**) which is a type of RNN and transformer based model which is A Lite version of Bidirectional Encoder Representation from Transformers

(ALBERT) ([Lan et al., 2019](#)). Additionally, I examine how various feature extraction methods, including TF-IDF and Word2Vec on ML algorithms, not DL algorithms..

By addressing these questions and finding the best performing algorithm along with feature extraction methods, this project seeks to contribute the development of the model that has large amount of accuracy and that can adapt to threats.

2 Related Work

([Sadia et al., 2023](#)) compared multiple machine learning algorithms for spam detection focusing on the tweets and they employed multiple classification techniques like Naive bias, KNN, SVM, and more. This paper emphasizes the efficiency of content-based feature extraction and importance of algorithm selection. ([Teja Nallamothu and Shais Khan, 2023](#)) explored various ML techniques for spam detection. Their study evaluates that SVM particularly well with an accurate feature extraction for effective spam filtering. ([Ghosh and Senthilrajan, 2023](#)) conducted a comparative analysis of ML techniques. Their work evaluates various models across different datasets focusing on email spam. The study highlights that SVM and Random Forest consistently deliver Superior performance. ([González-Carvajal and Garrido-Merchán, 2020](#)) compared BERT, the transformer based model with traditional ML techniques for text classification. Their study highlights the superior performance in capturing contextual information outperforming SVM across various datasets. ([Sahmoud and Mikki, 2022](#)) their research demonstrated that BERT effectively captures contextual nuances in text, outperforming traditional ML models on various datasets. It highlights the BERT's ability to improve accuracy, especially for more complex language patterns in spam messages. ([Yaseen et al., 2021](#)) explored models such as CNN and LSTM. The study

081 demonstrates that in particular LSTM provides a
 082 significant performance boost in identifying spam,
 083 especially in the large datasets. The paper empha-
 084 sizes the advantages of deep learning over tradi-
 085 tional methods in spam detection.

086 3 Data

087 This dataset for this project is acquired from
 088 Hugging-Face website ([Hugging Face Dataset](#))
 089 which comprises the spam emails, subject of the
 090 email, and label(spam, ham). The data set is speci-
 091 fically designed for spam classification. There are
 092 31,700 rows for training, and 2000 for testing. But,
 093 I joined all the data into single data-frame for better
 094 control over training, validation, and testing of the
 095 data.

096 Example: "*new clalls softtabs = instant rockhard*
 097 *erections simply dissolve half a pll under your*
 098 *tongue 10 min before action , for results that last*
 099 *all weekend . normal retail is 19 / pll order from*
 100 *us today at the price of 3 . 67 not interested (0 pt -*
 101 *0 ut).*" : spam

102 4 Methodology

103 4.1 Overview of Approach

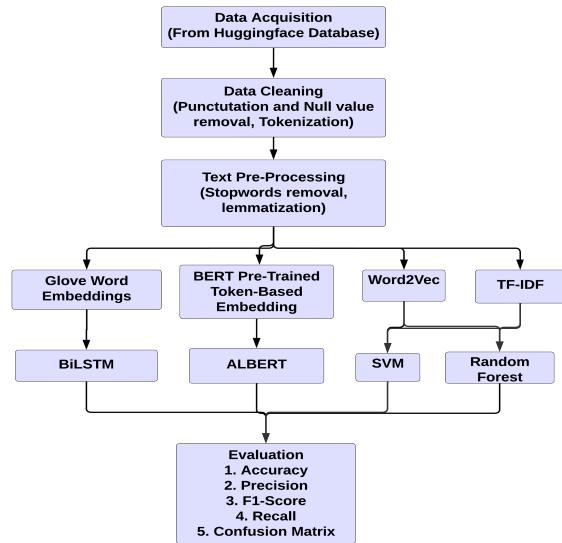


Figure 1: Overview of approach

104 4.2 Libraries and Tools

105 In this project, I have made use of various tools
 106 and libraries for different algorithms. For cleaning
 107 and pre-processing the data: **Pandas**, **Numpy**, and
 108 **Regex** are used. To normalize the text, **NLTK** li-
 109 brary used with tools like **Tokenizer**, **stop-words**,

110 and **Lemmatizer** are used. For vectorization of the
 111 text: **scikit-learn** library and **TF-IDF Vectorizer**,
 112 **Gensim** library : **Word2Vec** are used. For optimiz-
 113 ing the hyper-parameters for all the algorithms **Op-**
 114 **tuna** library is used. For implementation of **SVM**
 115 and **Random Forest**: **Scikit-Learn** is used. For
 116 the implementation of **ALBERT** and **BiLSTM**:
 117 **PyTorch**, **Transformers from Hugging-face (AL-**
 118 **BERT Model Documentation)** libraries are used.
 119 For accelerating the training of the deep learning
 120 models, **GPU(RTX 3050 Ti)** is leveraged by use
 121 of **NVIDIA CUDA**.

122 4.3 Dataset and Pre-Processing

123 As this project goal is spam detection, the dataset
 124 analyzed is "Enron Spam Email Dataset". It has
 125 7 columns, but for this project two of them are
 126 needed, i.e, "label", and "text". The "text" col-
 127 umn is concatenation of subject and message of the
 128 email. The dataset is balanced with spam:ham with
 129 50.8% : 49.2% of dataset. There are few missing
 130 missing values, the rows are dropped. The data
 131 is further processed by lower-casing, removal of
 132 punctuations, removal of numbers and special char-
 133 acters, removal of white spaces, tokenization, stop-
 134 words removal, and lemmatization of words. As
 135 TF-IDF cannot process tokens of text, text is joined
 136 as a sentence and stored in a new column. As it is
 137 spam detection the numbers, special characters and
 138 punctuation hold a lot of semantic meaning for a
 139 spam text, in section 5.1, experiments are conducted
 140 with and without removal of these elements.

141 4.4 Feature Representation

142 There were a total of 4 types of word embedding
 143 methods are used, with different parameters which
 144 can compliment the algorithm used.

145 4.4.1 TF-IDF

146 The Term Frequency - Inverse Document Fre-
 147 quency vectorizer transforms the data into nu-
 148 mercial features based on the importance of the
 149 words. More importance to frequent words, and
 150 less importance to common words. For the data
 151 after optimizing for the best parameters, it is im-
 152 plemented with max_features=12000 (vocabulary
 153 size), ngram_range=(1, 2).

154 4.4.2 Word2Vec

155 It is a neural-network based model that converts
 156 the words into vectors representation of many di-
 157 mensions such that semantic meaning of the word

158
159
160
161
162
is maintained which places alike words close in
multi-dimensional space. For the data, after op-
timizing for the best parameters, these are used
vector_size=300 , window=7, min_count=3, work-
ers=7, sg=1 (skip-gram).

163
164
**4.4.3 ALBERT: Pre-Trained Token-Based
Embedding**

165
166
167
168
169
170
171
The ALBERT tokenizer breakdown the text into to-
kens and convert them into vectors with numerical
ID and map them to a unique integer from a pre-
trained vocabulary. This is a type of embedding-
based vectorization. The text is vectorized into 128
dimensions. Padding and Truncation is performed
on each token.

172
173
**4.4.4 BiLSTM: GLOve Embeddings (Glove
Embeddings)**

174
175
176
177
178
179
180
181
182
GLOve is an unsupervised learning algorithm for
obtaining vector representations for words. Train-
ing is performed on aggregated global word-word
co-occurrence statistics from a corpus, and the re-
sulting representations showcase interesting linear
substructures of the word vector space. Pre-trained
GLOve embeddings of 300 dimensions are used in
which embedding layer converts input tokens into
dense vectors.

183
4.5 Model Architecture

184
4.5.1 Support Vector Machine

185
186
187
188
189
190
191
192
193
194
195
196
Support Vector Machine (SVM) is a powerful su-
pervised learning model used commonly for classi-
fications. In this project, SVM in optimized using
Optuna for both TF-IDF and Word2Vec datasets.
The svm model's hyper-parameters are regularization
parameter "C" (0.1 to 10), kernel type (lin-
ear, rbf, or poly), and polynomial degree (2 to 4),
which are all optimized to maximize the F1-score.
The model uses cache_size=10240 (10GB) and
max_iter =2000 for efficient and reproducible train-
ing. StandarScalar is applied to dense Word2Vec
data for consistency.

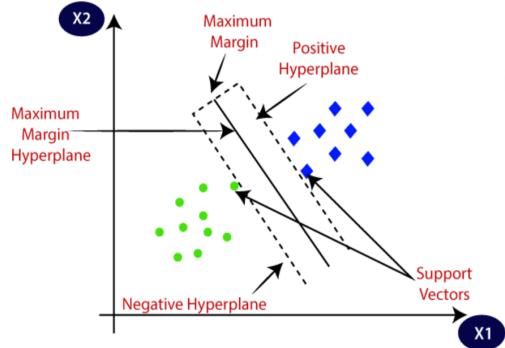


Figure 2: Support Vector Machine

197
4.5.2 Random Forest

198
199
200
201
202
203
204
205
206
207
208
209
210
211
212
213
214
215
216
217
218
219
220
Random Forest (RF) is a robust ensemble learn-
ing method that is used for classification tasks,
which combines multiple decision trees to improve
predictive accuracy and control over-fitting. In
this project, like SVM, RF is also optimized us-
ing Optuna to maximize the F1-score for both
datasets. Hyper-parameters such as, number
of trees: n_estimators= (100 to 300) with step
of 10, maximum depth of tree: max_depth=
(None,10,20,30,40,50), the minimum samples re-
quired to split an internal node: min_sample_split=
(2 to 10) are tuned for better model. The model
is trained using n_jobs=-1, which uses all the re-
sources of computer for parallel processing.

To avoid over-fitting for both SVM, and RF, the
training set is split into training and validation sub-
sets, with Optuna tuning based on the validation
set while leaving the test set for final evaluation.
5-Cross-validation is used during the optimization
to ensure robust hyper-parameter selection across
the multiple splits of the data. After 50 trials, the
best configurations for the algorithms are used to
train on both vectorized data (TF-IDF, Word2Vec).

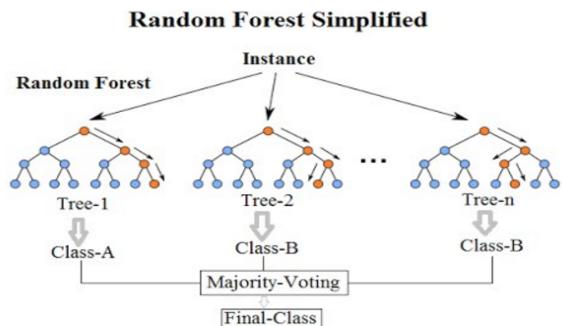


Figure 3: Random Forest

4.5.3 ALBERT (Lan et al., 2019)

ALBERT is a transformers model pre-trained on a large corpus of English data in a self-supervised fashion, the name of the model is "albert-base-v2". It has repeating layers which have same weights, it has small memory footprint but, has same computational cost as it has BERT like architecture. It has 12 repeating layers, 12 embedding dimension, 768 hidden dimension, 12 attention heads, and 11M parameters. This model is the fine tuned for spam classification task. For this project, the "Albert-Tokenizer" inputs maximum length of 128 tokens. Batching is done with batch_size=20 as the GPU cannot compute more. The model is trained for 6 epochs, optimized by "AdamW" with a learning rate of "5e-5" and weight decay of "1e-5" to prevent over-fitting. Gradient clipping (max 1.0) ensures stable updates, while StepLR scheduler reduces lr by 90% for every 2 epochs. Validation occurs at the end of each epoch, computing accuracy. Training of the model leverages GPU acceleration, with memory cleared afterwards to handle model's size in GPU.

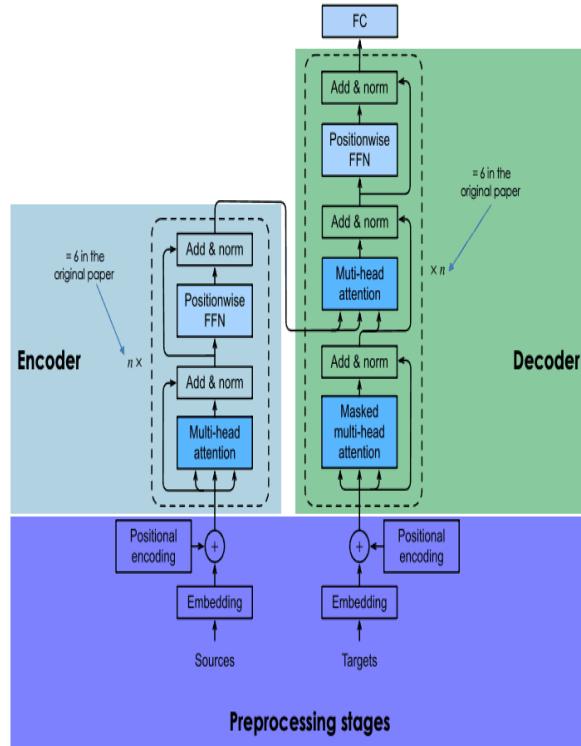


Figure 4: BERT

4.5.4 BiLSTM (Pytorch LSTM)

A Bidirectional LSTM(Long-Short-Term-Memory) is type of Recurring Neural Network known for their ability in capturing sequential data by processing

the input data in both forward and reverse direction. This method helps to capture patterns that might be missed when the data hidden layer only has the knowledge about past data by also passing the data from future. In this project, the data is sent to embedding layer to capture semantic relationships by leveraging GloVe embeddings (300 dimensions). The hyperparameters are number of layers stacked are: num_layers=2, hidden neurons in each layers are hidden_size=128, batch_size=20, learning_rate= 5e-5 and initial training epochs =10 with early stopping to prevent over-fitting. The AdamW optimizer enhanced by weight_decay= 5e-5 ensure effective parameter updates. There is dropout layer with rate=0.5, and early stopping which halts if validation loss doesn't improve for three consecutive epochs to avoid over-fitting. Validation occurs at the end of each epoch with cross-entropy loss function for performance on a separate validation dataset. This model leverages GPU acceleration for a faster processing.

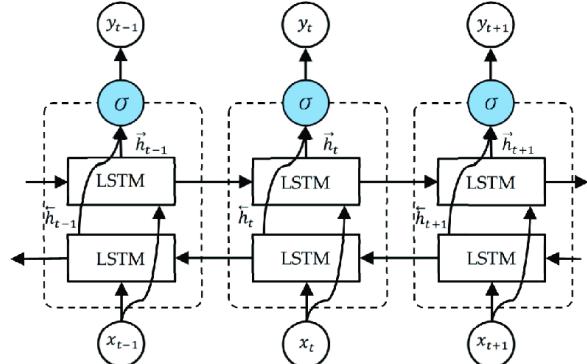


Figure 5: Bidirectional LSTM

4.5.5 Evaluation Metrics

The performance of the models are evaluated on the basis of different evaluation metrics like, **Accuracy**: it represents overall correctness of the model; by showing percentage of correctly predicted instances. **Precision**: it tells the number of instances of spam or ham are correctly identified as spam or ham. **Recall**: it tells, how effectively the model captures actual spam or ham messages. **F1-Score**: it balances precision and recall, making it very important as there is a need to balance the avoiding false positive with true positive. **Confusion Matrix**: it gives more detailed view of the model with TP, FP, TN, and FN states. By which the detailed breakdown of the model can be possible.

| | | POSITIVE | NEGATIVE | |
|---------------|----------|----------|----------|--|
| ACTUAL VALUES | POSITIVE | TP | FN | $Precision = \frac{TP}{TP + FP}$ |
| | NEGATIVE | FP | TN | $Recall = \frac{TP}{TP + FN}$ |
| | | | | $Accuracy = \frac{TP + TN}{TP + FP + FN + TN}$ |
| | | | | $F1 Score = 2 \times \frac{Precision \times Recall}{Precision + Recall}$ |

Figure 6: Evaluation Metrics

284 5 Experiments and Results

285 5.1 Experiments

286 As, the goal the project is to compare the algo-
 287 rithms, several experiments are conducted on the
 288 pre-processing steps, vectorization methods, and
 289 hyper-parameters.

290 5.1.1 Pre-Processing

291 The raw data is cleaned of null values and before
 292 the stop words removal, the pre-processing of data
 293 such as: **lower-casing, removal of numbers, re-**
 294 **moval of punctuations, and removal of empty**
 295 **space** in the text are performed. But, as it is spam
 296 detection, the punctuations, numbers, and empty
 297 space between the words have a lot of semantic
 298 meaning in determining a text spam or ham. So,
 299 the experiments are conducted twice such that all
 300 the algorithms are evaluated **with and without pre-**
 301 **processing** that are mentioned above.

302 5.1.2 Vectorization

303 In the project a total of four vectorization methods
 304 are used. To optimize the hyper parameters for TF-
 305 IDF, and Word2Vec, Optuna is used which gives
 306 the hyper parameters that performs best on these
 307 vectorization methods. For training the algorithms,
 308 the data that is fed is vectorized using these par-
 309 ameters. As ALBERT has its own text embeddings,
 310 and expects a certain dimension of vector. So, no
 311 experiments are conducted here. For BiLSTM, the
 312 embedding layer takes the vector from GlOve word
 313 embeddings, the experiments are conducted with
 314 100d, 200d, and 300d (d = dimensions) of GlOve
 315 embeddings. Finally, 300d is giving the best per-
 316 formance of the model.

317 5.1.3 Hyper-Parameters

318 As, discussed in section 4.5, all the hyper-
 319 parameters of the ML models are optimized by
 320 using Optuna. As, deep learning models takes enor-
 321 mous time and computation to optimize model with
 322 Optuna, the parameters like learning rate, weight

decay, optimizer, number of layers, batch size, and
 323 epochs are optimized manually by changing the
 324 values after each training cycle based on the evalua-
 325 tion of the models.

326 5.2 Results

327 5.2.1 Without Pre-Processing

329 1. Support Vector Machine (SVM)

```
Best parameters for TF-IDF: {'C': 3.4856532501539697, 'kernel': 'rbf'}
Best F1 Score for TF-IDF: 0.9936027915091596
Best parameters for Word2Vec: {'C': 5.308823214093609, 'kernel': 'rbf'}
Best F1 Score for Word2Vec: 0.992721979621543
[REDACTED]
TF-IDF SVM Test Accuracy: 0.9935
TF-IDF SVM Test Precision: 0.9910
TF-IDF SVM Test Recall: 0.9962
TF-IDF SVM Test F1 Score: 0.9936
TF-IDF SVM Test Confusion Matrix:
[[3283 31]
 [ 13 3417]]
C:\Users\watsa\AppData\Roaming\Python\Py
warnings.warn(
Word2Vec SVM Test Accuracy: 0.9926
Word2Vec SVM Test Precision: 0.9913
Word2Vec SVM Test Recall: 0.9942
Word2Vec SVM Test F1 Score: 0.9927
Word2Vec SVM Test Confusion Matrix:
[[3284 30]
 [ 20 3410]]
```

Figure 7: SVM-Without Pre-processing

330 2. Random Forest (RF)

```
Best parameters for TF-IDF: {'n_estimators': 300, 'max_depth': None, 'min_samples_split': 2}
Best F1 Score for TF-IDF: 0.9879901606135146
Best parameters for Word2Vec: {'n_estimators': 260, 'max_depth': 40, 'min_samples_split': 2}
Best F1 Score for Word2Vec: 0.988399071925754
TF-IDF Random Forest Test Accuracy: 0.9877
TF-IDF Random Forest Test Precision: 0.9808
TF-IDF Random Forest Test Recall: 0.9953
TF-IDF Random Forest Test F1 Score: 0.9880
TF-IDF Random Forest Test Confusion Matrix:
[[3247 67]
 [ 16 3414]]
Word2Vec Random Forest Test Accuracy: 0.9881
Word2Vec Random Forest Test Precision: 0.9833
Word2Vec Random Forest Test Recall: 0.9936
Word2Vec Random Forest Test F1 Score: 0.9884
Word2Vec Random Forest Test Confusion Matrix:
[[3256 58]
 [ 22 3408]]
```

Figure 8: RF-Without Pre-processing

331 3. ALBERT

| Classification Report: | | | | |
|-------------------------|-----------|--------|----------|---------|
| | precision | recall | f1-score | support |
| 0 | 0.99 | 0.98 | 0.99 | 3287 |
| 1 | 0.98 | 0.99 | 0.99 | 3446 |
| accuracy | | | 0.99 | 6733 |
| macro avg | 0.99 | 0.99 | 0.99 | 6733 |
| weighted avg | 0.99 | 0.99 | 0.99 | 6733 |
| Confusion Matrix: | | | | |
| [[3227 60] [22 3424]] | | | | |

Figure 9: ALBERT-Without Pre-processing

4. BiLSTM

| | precision | recall | f1-score | support |
|--------------------------|-----------|--------|----------|---------|
| 0 | 1.00 | 0.98 | 0.99 | 2510 |
| 1 | 0.98 | 1.00 | 0.99 | 2540 |
| accuracy | | | 0.99 | 5050 |
| macro avg | 0.99 | 0.99 | 0.99 | 5050 |
| weighted avg | 0.99 | 0.99 | 0.99 | 5050 |
| [[2466 44] [6 2534]] | | | | |

Figure 10: BiLSTM -Without Pre-processing

5. Comparison of all algorithms

| Algorithms/ Evaluation Metrics | Accuracy | Precision | Recall | F1-Score | Confusion Matrix |
|--------------------------------------|----------|-----------|--------|----------|----------------------|
| SVM (TF-IDF) | 99.35 | 99.10 | 99.62 | 99.36 | [3283 31 13 3417] |
| SVM (WORD2VEC) | 99.26 | 99.13 | 99.42 | 99.27 | [3284 30 20 3410] |
| Random Forest (TF-IDF) | 98.77 | 98.08 | 99.53 | 98.80 | [3247 67 16 3414] |
| Random Forest (WORD2VEC) | 98.81 | 98.33 | 99.36 | 98.84 | [3284 30 20 3410] |
| ALBERT | 98.78 | 98.28 | 99.6 | 98.82 | [3227 60 22 3424] |
| BILSTM | 99.00 | 98.29 | 99.76 | 99.02 | [2466 44 06 2534] |

Figure 11: All algorithms -Without Pre-processing

2. Random Forest (RF)

| |
|--|
| Best parameters for TF-IDF: {'n_estimators': 270, 'max_depth': None, 'min_samples_split': 4} |
| Best F1 Score for TF-IDF: 0.9874077290490665 |
| Best parameters for Word2Vec: {'n_estimators': 130, 'max_depth': 40, 'min_samples_split': 2} |
| Best F1 Score for Word2Vec: 0.9882455376578145 |
| TF-IDF Random Forest Test Accuracy: 0.9871 TF-IDF Random Forest Test Precision: 0.9805 TF-IDF Random Forest Test Recall: 0.9945 TF-IDF Random Forest Test F1 Score: 0.9874 TF-IDF Random Forest Test Confusion Matrix: [[3246 68] [19 3411]] Word2Vec Random Forest Test Accuracy: 0.9880 Word2Vec Random Forest Test Precision: 0.9838 Word2Vec Random Forest Test Recall: 0.9927 Word2Vec Random Forest Test F1 Score: 0.9882 Word2Vec Random Forest Test Confusion Matrix: [[3258 56] [25 3405]] |

Figure 13: RF-With Pre-processing

3. ALBERT

| Classification Report: | | | | |
|--|-----------|--------|----------|---------|
| | precision | recall | f1-score | support |
| 0 | 0.99 | 0.99 | 0.99 | 3332 |
| 1 | 0.99 | 0.99 | 0.99 | 3399 |
| accuracy | | | 0.99 | 6731 |
| macro avg | 0.99 | 0.99 | 0.99 | 6731 |
| weighted avg | 0.99 | 0.99 | 0.99 | 6731 |
| Confusion Matrix: [[3289 43] [34 3365]] | | | | |

Figure 14: ALBERT-With Pre-processing

5.2.2 With Pre-Processing

1. Support Vector Machine (SVM)

| |
|--|
| Best parameters for TF-IDF: {'C': 1.5432484576911387, 'kernel': 'rbf'} |
| Best F1 Score for TF-IDF: 0.9925839755707431 |
| Best parameters for Word2Vec: {'C': 1.6393469589175644, 'kernel': 'rbf'} |
| Best F1 Score for Word2Vec: 0.991723537098882 |
| TF-IDF SVM Test Accuracy: 0.9924 TF-IDF SVM Test Precision: 0.9901 TF-IDF SVM Test Recall: 0.9950 TF-IDF SVM Test F1 Score: 0.9926 TF-IDF SVM Test Confusion Matrix: [[3280 34] [17 3413]] C:\Users\vatsa\AppData\Roaming\Python\Pyth warnings.warn(Word2Vec SVM Test Accuracy: 0.9915 Word2Vec SVM Test Precision: 0.9879 Word2Vec SVM Test Recall: 0.9956 Word2Vec SVM Test F1 Score: 0.9917 Word2Vec SVM Test Confusion Matrix: [[3272 42] [15 3415]] |

Figure 12: SVM-With Pre-processing

| | precision | recall | f1-score | support |
|--------------------------|-----------|--------|----------|---------|
| 0 | 1.00 | 0.98 | 0.99 | 2522 |
| 1 | 0.98 | 1.00 | 0.99 | 2527 |
| accuracy | | | 0.99 | 5049 |
| macro avg | 0.99 | 0.99 | 0.99 | 5049 |
| weighted avg | 0.99 | 0.99 | 0.99 | 5049 |
| [[2481 41] [5 2522]] | | | | |

Figure 15: BiLSTM -With Pre-processing

5. Comparison of all algorithms

| Algorithms/ Evaluation Metrics | Accuracy | Precision | Recall | F1-Score | Confusion Matrix |
|--------------------------------------|----------|-----------|--------|----------|---------------------|
| SVM (TF-IDF) | 99.24 | 99.01 | 99.50 | 99.26 | [3280 34 173411] |
| SVM (WORD2VEC) | 99.15 | 98.79 | 99.56 | 99.17 | [3272 42 153415] |
| Random Forest (TF-IDF) | 98.71 | 98.05 | 99.45 | 98.74 | [3246 68 193411] |
| Random Forest (WORD2VEC) | 98.80 | 98.38 | 99.27 | 98.82 | [3258 56 253405] |
| ALBERT | 99.00 | 98.71 | 99.00 | 98.84 | [3289 43 343365] |
| BiLSTM | 99.09 | 98.40 | 99.80 | 99.06 | [2481 41 052522] |

Figure 16: All algorithms -With Pre-processing

5.3 Comparison of Models

In this section, I compare the performance of each model using the evaluation metrics, analyzing the strengths, and identifying best performing combinations.

5.3.1 Overall Model Performance

1. SVM (TF-IDF, Word2Vec): It shows the highest accuracy, and the F1-Score across both pre-processed and non-pre-processed dataset. The precision and recall remained high, demonstrating a robust performance. As, I tried to optimize the hyper-parameters, it takes a huge amount of time to train on and it cannot use all the resources in the computer as it cannot be defined in the implementation.

2. BiLSTM: It performed strongly as well, especially with pre-processing, achieving an F1-Score of 99.06% which indicates the capability to capture sequential information. However, as it is deep learning model with 2-layers stacked, even with GPU acceleration the training time is very heavy.

3. ALBERT: It achieved a commendable results across both the settings, performing close to Random Forest, but its computational requirements are very high, even-though I am using a vanilla model. If, I use a more complex model with more encoding and decoding layers, it may perform better than this setting. The algorithm works better with the preprocessed data than the raw data. The training time is highest among all the algorithms used.

4. Random Forest (TF-IDF, Word2Vec): It demonstrated slightly lower metrics than the other algorithms, but it was very efficient in finding the hyper-parameters for training the algorithm, As, it has implementation to use all the resources of computer. It also, offers balance between performance

and computational efficiency as it has least amount of training time in all the algorithms.

5.3.2 Best Model and Feature Extraction Combination

1. Top Performance: The SVM with TF-IDF combination yielded the best results in terms of accuracy(99.35) and F1-Score (99.36) without pre-processing. This indicates that SVM can leverage TF-IDF effectively for spam detection.

2. Sequential Data Model: The BiLSTM with pre-processing followed closely, excelling in recall(99.80), which is essential for spam detection tasks where minimizing False-Negatives is critical. Here, unlike ALBERT, BiLSTM captures context over the sequence using the LSTM units memory cells to retain the information across time steps. ALBERT perform extremely well on complex relationships across entire sequence, this dataset doesn't have enough complexity to leverage ALBERT. It would be reason for little decrease in performance.

5.3.3 Trade-Off Observations

1. SVM model trained faster and required fewer resources, ideal for the rapid iterations, with TF-IDF showing an edge over Word2Vec.

2. BiLSTM and **ALBERT**, while yielding competitive results, they demand a huge amount of computational power and slower to train due to sequential nature and stacked LSTM, and transformer-based approach of ALBERT. This trade-off makes them suitable for higher recall over training time and complex patterns in data.

3. Random Forest offered a good compromise of, balancing the training time and moderate consumption with a high accuracy. With more complex data, Random Forest may have potential for better performance.

5.3.4 Impact of Pre-Processing

1. Pre-processing improved BiLSTM recall and overall consistency for SVM and RF. In contrast, ALBERT saw a minimal increase of metrics, indicating the it can effectively handles unprocessed data.

2. Models with pre-processing showed a slight improvement in reducing False Positives, demonstrating that pre-processing can enhance model robustness, particularly for BiLSTM.

424 5.4 Note-Worthy Findings

425 Pre-processing steps improved model performance,
426 especially increasing the BiLSTM model's recall,
427 indicating its sensitivity to input quality. The high
428 accuracy of SVM with TF-IDF suggests its ef-
429 fectiveness in spam classification, aligning with
430 prior research. ALBERT performed competitively
431 but outperformed by SVM. RF struggled with
432 Word2Vec, likely due to sparse features affecting its
433 performance. As, the dataset is of 36,000 rows, it is
434 on the edge for better performance using ML or DL
435 algorithms for better results. Also, the data doesn't
436 have complex structures or features to leverage DL
437 algorithms. With more computational power, RF
438 and DL algorithms can train well by increasing lay-
439 ers, batch size, neurons, learning rate optimization,
440 and many more.

441 6 Discussion and Conclusion

442 This project aimed to find the best spam classi-
443 fication accuracy by evaluating the performance
444 of SVM, RF, ALBERT, and BiLSTM across
445 TF-IDF and Word2Vec feature extraction methods
446 on ML algorithms. SVM with TF-IDF with
447 pre-processing emerged as best-performing
448 combination vs Word2Vec. This suggests that
449 model and feature pairing significantly influence
450 the performance. Pre-processing the data improved
451 the performance significantly of all the algorithms,
452 but not much evidently in ALBERT which can be
453 inferred as ALBERT is immune to pre-processing
454 until certain extent. Dataset limitations include
455 potential noise, these can be mitigated with more
456 refined pre-processing. Ethical considerations in
457 deploying spam classification include potential
458 bias and importance of responsible data usage
459 to avoid increasing harmful stereotypes. Future
460 approach to ethical challenged should involve
461 model monitoring, periodic updates, and trans-
462 parent model decision process to ensure fairness
463 and reliability in the system. In conclusion, this
464 project underscore the significance of choosing
465 best model and feature combination for effective
466 spam classification. Future work might integrate
467 advanced deep learning architectures or ensemble
468 methods to further boost the performance and
469 robustness of automated spam detection systems.

471 References

- 472 ALBERT Model Documentation. Al-
473 bert model documentation. https://huggingface.co/docs/transformers/en/model_doc/albert#transformers.AlbertForSequenceClassification.
474
475 BILSTM. Lstm. <https://medium.com/@abhinav.mishra123/bidirectional-lstm-for-mnist-classification-using-pytorch>.
476
477 Argha Ghosh and A Senthilrajan. 2023. Comparison
478 of machine learning techniques for spam detection.
479 *Multimedia Tools and Applications*, 82(19):29227–
480 29254.
481
482 Glove Embeddings. Glove embeddings. <https://nlp.stanford.edu/projects/glove/>.
483
484 Santiago González-Carvajal and Eduardo C Garrido-
485 Merchán. 2020. Comparing bert against traditional
486 machine learning text classification. *arXiv preprint arXiv:2005.13012*.
487
488 Hugging Face Dataset. Spam dataset. https://huggingface.co/datasets/SetFit/enron_spam.
489
490 Zhenzhong Lan, Mingda Chen, Sebastian Goodman,
491 Kevin Gimpel, Piyush Sharma, and Radu Soricut.
492 2019. ALBERT: A lite BERT for self-
493 supervised learning of language representations.
494 *CoRR*, abs/1909.11942.
495
496 Pytorch LSTM. Pytorch lstm. <https://pytorch.org/docs/stable/generated/torch.nn.LSTM.html>.
497
498 Azeema Sadia, Fatima Bashir, Reema Qaiser Khan, and
499 A Khalid. 2023. Comparison of machine learning
500 algorithms for spam detection. *Journal of Advances in Information Technology*, 14(2):178–184.
501
502 Thaer Sahmoud and Dr Mohammad Mikki. 2022.
503 Spam detection using bert. *arXiv preprint arXiv:2206.02443*.
504
505 Phani Teja Nallamothu and Mohd Shais Khan. 2023.
506 Machine learning for spam detection. *Asian Journal of Advances in Research*, 6(1):167–179.
507
508 Qussai Yaseen et al. 2021. Spam email detection us-
509 ing deep learning techniques. *Procedia Computer Science*, 184:853–858.
510
511